

FYS-STK3155/4155 Applied Data Analysis and Machine Learning - Project 2: Classification and Regression

Lotsberg, Bernhard Nornes
Nguyen, Anh-Nguyet Lise

<https://github.com/liseanh/FYS-STK4155-project2/>

October - November 2019

Abstract

blip bloop opinion wrong

1 Introduction

Classification in statistical analysis is a useful tool, e.g. for predicting outcomes of various situations or classifying and sorting large amounts of data.

The aim of this project is to study classification and regression problems through our own implementation of logistic regression and a multilayer perceptron (MLP) in Python. The particular data set we will be studying for classification has been used in a prior research paper by Yeh, I. C. and Che-hui Lien about data mining techniques [5], which we will compare some of our results with. The full data set can be downloaded from the UCI Machine Learning Repository [4]. The data set contains certain features of credit card clients' default payment from a Taiwanese bank. For the regression problem we will use the MLP to approximate Franke's function and compare with results from a prior project where we used ordinary least squares, ridge and lasso regression to approximate it. [2].

2 Learning methods

2.1 Logistic Regression (LR)

Logistic regression (LR) is a statistical model that can be used to predict a binary dependent variable. It functions in a manner similar to linear regression (hence the name), the difference being that the outputs are put through an activation function. Usually this activation function is the Sigmoid function, which for a binary case with logistic regression is defined as

$$p(y = 1|X, \beta) = \frac{1}{1 + e^{-X\beta}} = \frac{e^{X\beta}}{e^{X\beta} + 1}. \quad (1)$$

Here X is the design matrix, and β is a vector containing the weights assigned to each feature in X . It is useful to use another cost function than the mean squared error, as the output from our logistic regression is defined as $\hat{y} = p(X) \in [0, 1]$, where we say that outcome 0 is predicted if $\hat{y} < 0.5$ and outcome 1 else. The cost function most commonly used in this case is called the cross entropy, and is

defined as

$$\begin{aligned} \mathcal{C}_{\text{LR}}(\beta) = & - \sum_{i=0}^{n-1} [y_i \log p(y_i|x_i, \beta) \\ & + (1 - y_i) \log(1 - p(y_i|x_i, \beta))] \end{aligned} \quad (2)$$

We find the weights β by solving $\operatorname{argmin}_{\beta}$. Like with LASSO regression which we used in our previous project, this has no analytical solution and must be estimated using gradient descent [2]. This along with its stochastic sibling is described in section 2.4.

In this project we will use LR on the Taiwanese credit card data to try to predict default ($y = 1$) or non-default ($y = 0$) payment and evaluate the model's predictive accuracy in this particular classification problem.

2.2 Neural Networks (NN)

An artificial neural network (NN) is a computational model consisting of interconnected nodes. The interconnected nodes aim to emulate a simplified biological neural network and neuronal firing in a brain, and are therefore also commonly referred to as neurons. The node performs a weighted sum of its inputs that is subsequently passed through a mathematical function to determine its output. This mathematical function is called an activation function $f(z)$, and should emulate neuronal firing.

There is a wide variety of different neural networks. Commonly, they consist of layers of nodes separated into the input layer and output layer, and may also contain one or more in-between layers called hidden layers. One such NN is the multilayer perceptron (MLP), which is what we will be using in this project for both classification and regression.

2.3 Multilayer perceptron (MLP)

2.3.1 Feed-forward

The multilayer perceptron is a feed-forward neural network (FFNN), which means that the information flows forward only, starting from the input layer and to the output layer. Additionally, if each of the nodes in a layer is connected to all of the nodes in the succeeding layer, the network is fully connected. The inputs of the node are the weighted outputs of the nodes from the preceding layer, in addition to a bias term that can control whether or not the neuron fires if all the inputs are zero [3]. The weighted sum of the inputs of each node is called the activation.

Mathematical algorithm

Input layer

Starting with the input layer, which is the first layer in the MLP, the activation is calculated using the input coordinates x_j ,

$$z_i^1 = \sum_{j=1}^{M_1} w_{ij}^1 x_j + b_i^1, \quad (3)$$

where the superscript 1 indicates the first layer, M_1 is the number of inputs to the i th node in the first layer, b_i is the bias and w_{ij} represents the weights.

The output of the nodes in the input layer is determined by the activation function $f(z)$,

$$f(z_i^1) = f \left(\sum_{j=1}^{M_1} w_{ij}^1 x_j + b_i^1 \right) \quad (4)$$

Hidden layers and output layer

Similarly for the subsequent layers; the hidden layers and the output layer, the activation of the j th neuron of layer l is defined as

$$z_j^l = \sum_{i=1}^{M_{l-1}} w_{ij}^l a_j^{l-1} + b_j^l, \quad (5)$$

where b_j^l and w_{ij}^l are the biases and weights at layer l , M_{l-1} is the number of nodes at layer $l-1$ and $a_j^{l-1} = f(z_j^{l-1})$. The output of each node is then decided by passing the activation through the activation function,

$$f(z_j^l) = f \left(\sum_{i=1}^{M_{l-1}} w_{ij}^l a_i^{l-1} + b_j^l \right) \quad (6)$$

This is nearly identical to the method for the input layer, except that the inputs of these layers are the outputs from the previous layer.

2.3.2 Backpropagation

In order for the NN to learn, the weights and biases are initialized with values that we will discuss shortly. The weights and biases are then optimized to minimize the cost function through a process called backpropagation, where we iterate backwards from the last layer to the first hidden layer. Another feed-forward process is initiated from the input layer to the output layer with the new biases and weights. If the cost function is not yet sufficiently minimized, then backpropagation is performed again. This process is repeated until the cost function is optimized.

Mathematical algorithm

To calculate the optimal biases and weights for the problem, we initialize the gradients of the cost function \mathcal{C} with respect to the weights W and biases b at the output layer $l = L$ and the output error δ_L as

$$\frac{\partial \mathcal{C}}{\partial w_{jk}^L} = \delta_j^L a_k^{L-1}, \quad (7)$$

$$\frac{\partial \mathcal{C}}{\partial b_j^L} = \delta_j^L, \quad (8)$$

$$\delta_j^L = f'(z_j^L) \frac{\partial \mathcal{C}}{\partial a_j^L}, \quad (9)$$

before propagating backwards through the hidden layers using the general equations

$$\frac{\partial \mathcal{C}}{\partial w_{jk}^l} = \delta_j^l a_k^{l-1}, \quad (10)$$

$$\frac{\partial \mathcal{C}}{\partial b_j^l} = \delta_j^l, \quad (11)$$

$$\delta_j^l = \sum_k \delta_k^{l+1} w_{kj}^{l+1} f'(z_j^l). \quad (12)$$

The full derivation of these equations can be found in *Neural networks, from the simple perceptron to deep learning* by Morten Hjorth-Jensen [1].

2.3.3 Choosing activation function

There are certain properties that an activation function $f(z)$ should possess [3]:

1. $f(z)$ should be differentiable (continuous).
2. $f(z)$ should be non-constant and reach saturation at both ends of the range.
3. $f(z)$ should change quickly between the saturation values at the middle of the range.

It is clear from Equation (12) in the backpropagation algorithm that the activation function $f(z)$ should be differentiable, while the remaining properties are mainly related to the emulation of firing of neurons in a brain.

There are several functions that meet these requirements. Historically, the sigmoid function has been a popular choice of activation function until recently. To compare our results with the original paper from 2009 by I-Cheng Yeh and Che-hui Lien [5], we have chosen to use the sigmoid function in our implementation of the MLP, as it is likely the same function they used. The function is given by

$$f(z) = \frac{1}{1 + e^{-z}}. \quad (13)$$

Its gradient is given by

$$\frac{\partial f}{\partial z} = \frac{e^{-z}}{(1 + e^{-z})^2}. \quad (14)$$

However, there are cases where we do not want a sigmoidal output activation function, such as for regression problems where the outputs should be from a continuous range instead of binary. Instead of a sigmoidal activation function for the regression case at output layer $l = L$, we implement a linear activation function, such that

$$f(z_j^L) = z_j^L = \sum_{i=1}^{M_{L-1}} w_{ij}^L a_j^{L-1} + b_j^L. \quad (15)$$

For our binary classification case, we keep the sigmoid function as the output activation function.

2.3.4 Choosing cost function

Looking at the equations for backpropagation, it is clear that the chosen cost function \mathcal{C} should be differentiable as well. Additionally, the cost function should be chosen with the desired functionality in mind, i.e. the cost functions for regression and classification should be different.

For the classification part of this project we have chosen to use the logistic loss as cost function, which is the negative log-likelihood,

$$\begin{aligned} \mathcal{C}_{\text{NN}}^C(\mathbf{W}, \mathbf{b}) = & - \sum_{i=0}^{n-1} \left[a_i^l \log p(a_i^l | z_i^l, \mathbf{W}, \mathbf{b}) \right. \\ & \left. + (1 - a_i^l) \log (1 - p(a_i^l | z_i^l, \mathbf{W}, \mathbf{b})) \right], \end{aligned} \quad (16)$$

with $a_i^l = f(z_i^l)$ as before. The gradient at the last layer is

$$\frac{\partial \mathcal{C}_{\text{NN}}^C}{\partial a_i^L} = a_i^L - t_i, \quad (17)$$

where t_i is the target.

For the regression part we will use the cost function

$$\mathcal{C}_{\text{NN}}^R = \frac{1}{2} \sum_{i=0}^{n-1} (t_i - a_i^L)^2. \quad (18)$$

The gradient in the last layer is given by

$$\frac{\partial \mathcal{C}_{\text{NN}}^R}{\partial a_i^L} = t_i - a_i^L. \quad (19)$$

The gradients in Equation (17) and (19) can then be inserted into Equation (12) in the backpropagation algorithm.

2.3.5 Initialising the weights and biases

The biases can be initialized to zero, but we have chosen an initial value of 0.01 to ensure that all of the neurons have some initial output. Initializing the weights to zero, however, will result in all neurons outputting the same value. Instead, the weights are initialized with values drawn from a uniform distribution such that $w_{kj} \in (-1/\sqrt{n}, 1/\sqrt{n})$, where n is the amount of nodes in the input layer, to ensure uniform learning [3].

2.3.6 Regularisation

2.4 Stochastic Gradient Descent (SGD)

It is clear that in order to find the best possible fit, we need to optimize the cost function \mathcal{C} by finding its minimum. A common method to achieve this is the gradient descent (GD) method, in which the parameters θ are iteratively adjusted in the direction of the largest negative value of the gradient for a given number of iterations or until it reaches a given tolerance. Mathematically, this is expressed as

$$\theta_{i+1} = \theta_i - \eta \nabla \mathcal{C}(\theta_i), \quad (20)$$

where η is the learning rate, which is a hyperparameter that controls the step length and by extension the convergence time. For smaller values of η , the method will take longer to converge or might not converge at all within a desired time frame. For larger values of η , the method might be unstable or pass the minimum altogether and diverge. The parameter θ is β in the LR case, and the weights W and biases b in the MLP case.

However, calculating the gradient on the entire data set can be computationally expensive and inefficient for large amounts of data. Additionally, there is a high possibility of a local minimum being misinterpreted as a global minimum by the algorithm. To alleviate these problems, we can use stochastic gradient descent (SGD) with minibatches. A minibatch is a subset of the data, on which we can perform GD. By using stochasticity to perform gradient descent on randomly chosen minibatches of size M , we have a more efficient way to approximate the gradient of the total data set as it might not need to use the entire set. Additionally, the stochasticity reduces the possibility of getting stuck in a local minimum.

2.5 Choice of hyperparameters

3 Data

3.1 Classification - Credit card client data

For the classification part of this project, we are using credit card payment data from a Taiwanese bank downloaded from the UCI Machine Learning Repository. The response variable is a binary variable of default payment with Yes = 1, No = 0. The original data set consists of 30 000 observations, with X amount of observations with default payments. There are 23 explanatory variables, cited from the original paper they are de-

scribed as [5]:

- X1: Amount of the given credit (NT dollar): it includes both the individual consumer credit and his/her family (supplementary) credit.
- X2: Gender (1 = male; 2 = female).
- X3: Education (1 = graduate school; 2 = university; 3 = high school; 4 = others).
- X4: Marital status (1 = married; 2 = single; 3 = others).
- X5: Age (year).
- X6 - X11: History of past payment. We tracked the past monthly payment records (from April to September, 2005) as follows: X6 = the repayment status in September, 2005; X7 = the repayment status in August, 2005; . . . ; X11 = the repayment status in April, 2005. The measurement scale for the repayment status is: -1 = pay duly; 1 = payment delay for one month; 2 = payment delay for two months; . . . ; 8 = payment delay for eight months; 9 = payment delay for nine months and above.
- X12-X17: Amount of bill statement (NT dollar). X12 = amount of bill statement in September, 2005; X13 = amount of bill statement in August, 2005; . . . ; X17 = amount of bill statement in April, 2005.
- X18-X23: Amount of previous payment (NT dollar). X18 = amount paid in September, 2005; X19 = amount paid in August, 2005; . . . ; X23 = amount paid in April, 2005.

3.2 Regression - Franke's function

For the regression part of this project we will be performing regression analysis on Franke's function $g(x, y)$ with added Gaussian noise $\varepsilon \sim N(0, \sigma^2)$. Franke's function

is given by

$$\begin{aligned}
g(x, y) = & \frac{3}{4} \exp \left(-\frac{(9x-2)^2}{4} - \frac{(9y-2)^2}{4} \right) \\
& + \frac{3}{4} \exp \left(-\frac{(9x+1)^2}{49} - \frac{(9y+1)^2}{10} \right) \\
& + \frac{1}{2} \exp \left(-\frac{(9x-7)^2}{4} - \frac{(9y-3)^2}{4} \right) \\
& - \frac{1}{5} \exp \left(-(9x-4)^2 - (9y-7)^2 \right)
\end{aligned} \tag{21}$$

and is defined on $x, y \in [0, 1]$. The data we are fitting is given by

$$G(\mathbf{x}, \mathbf{y}) = g(\mathbf{x}, \mathbf{y}) + \varepsilon,$$

where \mathbf{x}, \mathbf{y} are vectors of uniformly spaced values from 0 and 1 of length n_x and n_y respectively. To directly compare with our previous project, *Project 1: Regression analysis and resampling methods* [2], we choose to generate and analyze the data sets `franke0`, `franke1`, `franke2` as configured in Table 1. To generate the points, we grid over n_x points in x -direction and n_y points in y -direction.

WE NEED TO TALK MORE ABOUT OUR PREVIOUS ANALYSIS AND WHAT KIND OF METHODS AND PARAMETERS WE USED

Table 1: Table of the configurations for our generated data sets using Franke's function given in Equation (21) with added noise $\varepsilon \sim N(0, \sigma^2)$. n_x and n_y indicate the number of points used to produce the grid in their respective directions.

Data set	Data points	n_x	n_y	σ
<code>franke0</code>	400	20	20	1.0
<code>franke1</code>	400	20	20	0.10
<code>franke2</code>	40 000	200	200	0.10

4 Model evaluation

4.1 Regression

To evaluate the performance of our regression model, we consider the R^2 score, given by

$$R^2(\mathbf{y}, \hat{\mathbf{y}}) = 1 - \frac{\sum_{i=0}^{n-1} (y_i - \hat{y}_i)^2}{\sum_{i=0}^{n-1} (y_i - \bar{y})^2}, \tag{22}$$

where \mathbf{y} is the given data, $\hat{\mathbf{y}}$ is the model and \bar{y} is the mean value of \mathbf{y} .

4.2 Classification

To evaluate the performance of our classification model, we consider the accuracy score, given by

$$\text{accuracy} = \frac{\sum_{i=1}^n I(t_i = y_i)}{n}, \tag{23}$$

where t_i is the target, y_i is the model output, n is the number of samples and I is the indicator function,

$$I = \begin{cases} 1, & t_i = y_i \\ 0, & t_i \neq y_i \end{cases}.$$

5 Results




Figure 1: Accuracy score of validation set of the credit card data using the logistic regression classifier. The values of the shrinkage parameter λ were chosen using randomized search.

Table 2: Table of error rates and area ratios for classification of the credit card data using logistic regression (LR) and the multilayer perceptron (MLP) neural network (NN).

Method	Error rate	Area ratio default	Area ratio non-default
LR	0.31	0.43	0.12
NN	0.22	0.55	0.16

Table 3: Table of R² scores for our neural network (NN) regression models on the three Franke data sets and the ordinary least squares (OLS) model from our previous study [2].

	R ²	franke0	franke1	franke2
NN	Train	0.0052	0.59	0.84
	Test	-0.057	0.59	0.83
OLS	Train	0.024	0.87	N/A
	Test	-0.040	0.82	N/A

Table 4: Table of the best hyperparameters values for shrinkage λ and learning rate η for the logistic regression (LR) and neural network (NN) models. The parameters were found using 5-fold cross validation. The minibatch size was kept constant at $M = 200$.

Model	Shrinkage λ	Learning rate η
LR credit data	N/A	$4.4 \cdot 10^{-4}$
NN credit data	$6.9 \cdot 10^{-7}$	$7.7 \cdot 10^{-2}$
NN franke0	$4.1 \cdot 10^{-6}$	$9.1 \cdot 10^{-5}$
NN franke1	$6.4 \cdot 10^{-8}$	$1.5 \cdot 10^{-4}$
NN franke2	$3.6 \cdot 10^{-10}$	$3.7 \cdot 10^{-4}$




Figure 2: Accuracy score of validation set of the credit card data using the multilayer perceptron classifier. The values of the shrinkage parameter λ and the learning rate η were chosen using randomized search.




Figure 4: idk yet




Figure 3: idk yet




Figure 5: R^2 score of validation set of the Franke function data set with 400 total points $n_x = n_y = 20$, $\sigma = 1.0$ using the multilayer perceptron regressor. The values of the shrinkage parameter λ and the learning rate η were chosen using randomized grid search.




Figure 6: R^2 score of validation set of the Franke function data set with 400 total points, $n_x = n_y = 20$, $\sigma = 0.1$ using the multilayer perceptron regressor. The values of the shrinkage parameter λ and the learning rate η were chosen using randomized grid search.




Figure 7: R^2 score of validation set of Franke function data with 40 000 total points, $n_x = n_y = 200$, $\sigma = 0.1$ using the multilayer perceptron regressor. The values of the shrinkage parameter λ and the learning rate η were chosen using randomized grid search.

6 Discussion

fill

7 Conclusion

fill

References

- [1] Morten Hjorth-Jensen. Neural networks, from simple perceptron to deep learning. <https://compphysics.github.io/MachineLearning/doc/pub/NeuralNet/pdf/NeuralNet-minted.pdf>, 2019. Retrieved: 07-11-2019.
- [2] Bernhard Nornes Lotsberg and Anh-Nguyet Lise Nguyen. *Project 1: Regression analysis and resampling methods*.




Figure 8: 3D plot of the Franke function with added noise $\sim \mathcal{N}(0, \sigma^2)$ with $\sigma = 1.0$. The dotted markers represent the regression model found using MLP. The left plot shows the model fitted on the training set, the right shows the model fitted on the test set. The total data set consist of 400 points, using two thirds ($2/3$) as the training set and the remaining points ($1/3$) as the test set.




Figure 9: 3D plot of the Franke function with added noise $\sim \mathcal{N}(0, \sigma^2)$ with $\sigma = 0.1$. The dotted markers represent the regression model found using MLP. The left plot shows the model fitted on the training set, the right shows the model fitted on the test set. The total data set consist of 400 points, using two thirds ($2/3$) as the training set and the remaining points ($1/3$) as the test set.




Figure 10: 3D plot of the Franke function with added noise $\sim \mathcal{N}(0, \sigma^2)$ with $\sigma = 0.1$. The dotted markers represent the regression model found using MLP. The left shows the model fitted on the training set, the right shows the model fitted on the test set. The total data set consist of 40 000 points, using two thirds ($2/3$) as the training set and the remaining points ($1/3$) as the test set.

FYS-STK3155/4155 Applied Data Analysis and Machine Learning. 2019.

- [3] Stephen Marsland. *Machine Learning: An Algorithmic Perspective, Second Edition*, pages 71–107. Chapman and Hall/CRC, 2014.
- [4] UCI Machine Learning Repository. Default of credit card clients Data Set. <https://archive.ics.uci.edu/ml/datasets/default+of+credit+card+clients>. Retrieved: 08-10-2019.
- [5] I-Cheng Yeh and Che-hui Lien. The comparisons of data mining techniques for the predictive accuracy of probability of default of credit card clients. *Expert Systems with Applications*, 36(2):2473–2480, 2009.